


RF-DETR: faster metric computation

Actual one-pass metric · Nano predictions · 5,000 COCO images · maxDets=500

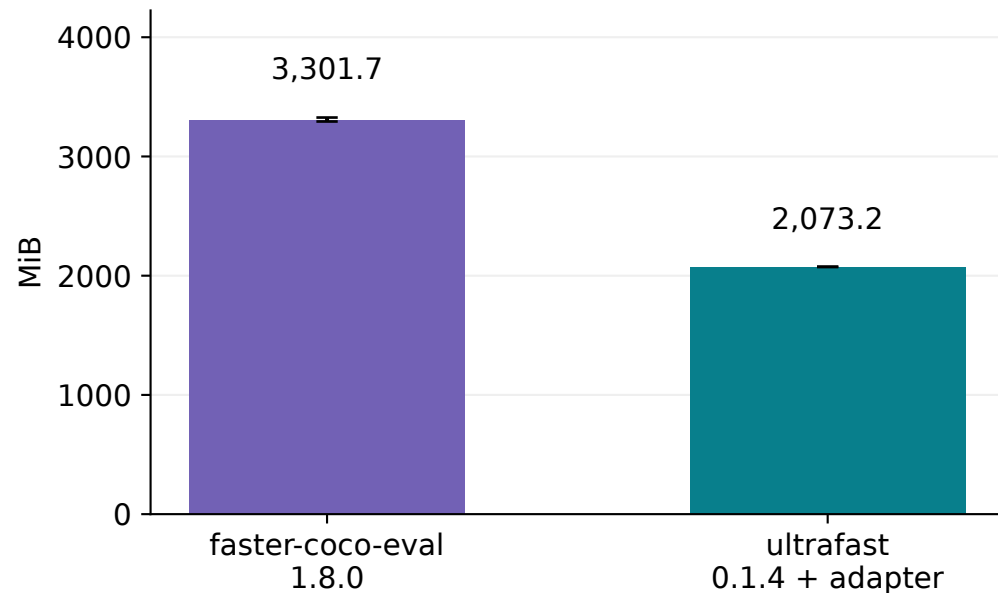
2.09× faster

Metric update + compute



37.2% less memory

Peak process memory



3 fresh processes per backend; median and min-max. Two CPU threads on a shared EPYC 9554.

Aggregate and per-class tensors match exactly. Metric replay excludes inference and trace loading.